

ASTEROIDES CON PELIGRO POTENCIAL PARA LA TIERRA

Recientemente, la NASA llevó a cabo el primer intento para desviar un asteroide que pudiera suponer un peligro de colisión con la Tierra (<https://www.nasa.gov/planetarydefense/dart/>).

Utilizando datos procedentes de la NASA (https://ssd.jpl.nasa.gov/tools/sbdb_query.html) intentaremos evaluar mediante modelos de machine learning si es posible predecir si un asteroide es potencialmente peligroso para nuestro planeta.

1. ANÁLISIS EXPLORATORIO DE LOS DATOS

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelBinarizer, LabelEncoder, OrdinalEncoder, MinM
from sklearn.model_selection import StratifiedShuffleSplit, train_test_split, GridSe
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report,
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import precision_recall_fscore_support as score

import warnings
warnings.filterwarnings('ignore', module='sklearn')
warnings.filterwarnings('ignore', module='IPython')
```

```
In [2]: filepath = 'C:/Users/NITROPC/Desktop/DATA SCIENCE/CERTIFICACION MACHINE LEARNING/06
data = pd.read_csv(filepath, sep = ',')
```

```
In [3]: data.head()
```

```
Out[3]:
```

	id	spkid	full_name	pdes	name	prefix	neo	pha	H	diameter	...	sigma_i
0	a0000001	2000001	1 Ceres	1	Ceres	NaN	N	N	3.40	939.400	...	4.608900e-09
1	a0000002	2000002	2 Pallas	2	Pallas	NaN	N	N	4.20	545.000	...	3.469400e-06
2	a0000003	2000003	3 Juno	3	Juno	NaN	N	N	5.33	246.596	...	3.223100e-06
3	a0000004	2000004	4 Vesta	4	Vesta	NaN	N	N	3.00	525.400	...	2.170600e-07
4	a0000005	2000005	5 Astraea	5	Astraea	NaN	N	N	6.90	106.699	...	2.740800e-06

5 rows × 45 columns



A continuación se recoge la descripción básica de las columnas, de acuerdo a lo indicado en la web JPL:

- SPK-ID: Object primary SPK-ID
- Object ID: Object internal database ID
- Object fullname: Object full name/designation
- pdes: Object primary designation
- name: Object IAU name
- NEO: Near-Earth Object (NEO) flag
- PHA: Potentially Hazardous Asteroid (PHA) flag
- H: Absolute magnitude parameter
- Diameter: object diameter (from equivalent sphere) km Unit
- Albedo: Geometric albedo
- Diameter_sigma: 1-sigma uncertainty in object diameter km Unit
- Orbit_id: Orbit solution ID
- Epoch: Epoch of osculation in modified Julian day form
- Equinox: Equinox of reference frame
- e: Eccentricity
- a: Semi-major axis au Unit
- q: perihelion distance au Unit
- i: inclination; angle with respect to x-y ecliptic plane
- tp: Time of perihelion passage TDB Unit
- moid_Id: Earth Minimum Orbit Intersection Distance au Unit

Nuestra variable objetivo es la variable PHA, que indica si un asteroide es potencialmente peligroso o no.

In [4]: `data.shape`

Out[4]: (958524, 45)

In [5]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 958524 entries, 0 to 958523
Data columns (total 45 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    958524 non-null object
1   spkid                 958524 non-null int64
2   full_name            958524 non-null object
3   pdes                  958524 non-null object
4   name                  22064 non-null object
5   prefix                18 non-null    object
6   neo                   958520 non-null object
7   pha                   938603 non-null object
8   H                     952261 non-null float64
9   diameter              136209 non-null float64
10  albedo                 135103 non-null float64
11  diameter_sigma        136081 non-null float64
12  orbit_id              958524 non-null object
13  epoch                 958524 non-null float64
14  epoch_mjd             958524 non-null int64
15  epoch_cal             958524 non-null float64
16  equinox               958524 non-null object
```

```
17 e          958524 non-null float64
18 a          958524 non-null float64
19 q          958524 non-null float64
20 i          958524 non-null float64
21 om         958524 non-null float64
22 w          958524 non-null float64
23 ma         958523 non-null float64
24 ad         958520 non-null float64
25 n          958524 non-null float64
26 tp         958524 non-null float64
27 tp_cal     958524 non-null float64
28 per        958520 non-null float64
29 per_y      958523 non-null float64
30 moid       938603 non-null float64
31 moid_ld    958397 non-null float64
32 sigma_e    938602 non-null float64
33 sigma_a    938602 non-null float64
34 sigma_q    938602 non-null float64
35 sigma_i    938602 non-null float64
36 sigma_om   938602 non-null float64
37 sigma_w    938602 non-null float64
38 sigma_ma   938602 non-null float64
39 sigma_ad   938598 non-null float64
40 sigma_n    938602 non-null float64
41 sigma_tp   938602 non-null float64
42 sigma_per  938598 non-null float64
43 class      958524 non-null object
44 rms        958522 non-null float64
dtypes: float64(33), int64(2), object(10)
memory usage: 329.1+ MB
```

Veamos nuestra variable objetivo.

```
In [6]: data['pha'].value_counts(normalize=True)
```

```
Out[6]: N    0.997799
        Y    0.002201
        Name: pha, dtype: float64
```

Como podemos ver, los datos de nuestra variable objetivo están muy poco balanceados.

Veamos cómo se distribuyen los valores dentro de cada una de las variables.

```
In [7]: pd.DataFrame([[i, len(data[i].unique())] for i in data.columns],
                    columns=['Variable', 'Unique Values']).set_index('Variable')
```

```
Out[7]:
```

Variable	Unique Values
id	958524
spkid	958524
full_name	958524
pdes	958524
name	22065
prefix	2
neo	3
pha	3
H	9490

Unique Values	
Variable	
diameter	16592
albedo	1058
diameter_sigma	3055
orbit_id	4690
epoch	5246
epoch_mjd	5246
epoch_cal	5246
equinox	1
e	958444
a	958509
q	958509
i	958414
om	958518
w	958519
ma	958520
ad	958506
n	958514
tp	958519
tp_cal	958499
per	958511
per_y	958512
moid	314301
moid_ld	314302
sigma_e	254741
sigma_a	273298
sigma_q	248139
sigma_i	215742
sigma_om	223156
sigma_w	262720
sigma_ma	266817
sigma_ad	269242
sigma_n	251751
sigma_tp	291247
sigma_per	282688
class	13

Unique Values	
Variable	
rms	64387

2. LIMPIEZA DE DATOS

Eliminemos algunas columnas que no aportarán ningún tipo de información.

```
In [8]: data1 = data.drop(['id', 'full_name', 'pdes', 'name', 'prefix', 'equinox'], axis='col')
```

Veamos a continuación los datos perdidos dentro de nuestra dataframe.

```
In [9]: num_missing = data.isnull().sum()
pctg_missing = data.isnull().sum().apply(lambda x: x/data.shape[0]*100)
```

```
In [10]: missing_data = pd.DataFrame({'Number of Missing': num_missing,
                                     'Percentage of Missing': pctg_missing})

missing_data['Percentage of Missing'].sort_values(ascending = False)
```

```
Out[10]: prefix          99.998122
name          97.698128
albedo        85.905100
diameter_sigma 85.803068
diameter       85.789714
sigma_ad        2.078821
sigma_per       2.078821
sigma_e         2.078404
sigma_a         2.078404
sigma_q         2.078404
sigma_i         2.078404
sigma_om        2.078404
sigma_w         2.078404
sigma_ma        2.078404
sigma_n         2.078404
sigma_tp        2.078404
pha             2.078300
moid           2.078300
H              0.653400
moid_ld        0.013250
per            0.000417
ad             0.000417
neo            0.000417
rms            0.000209
ma             0.000104
per_y          0.000104
class          0.000000
id             0.000000
tp_cal         0.000000
equinox        0.000000
full_name      0.000000
pdes           0.000000
orbit_id       0.000000
epoch          0.000000
epoch_mjd      0.000000
epoch_cal      0.000000
e              0.000000
tp             0.000000
a              0.000000
q              0.000000
i              0.000000
```

```

om          0.000000
spkid       0.000000
n           0.000000
w           0.000000
Name: Percentage of Missing, dtype: float64

```

Como podemos ver hay cinco variables con valores perdidos por encima del 80%. Estas variables no pueden ser calculadas ni sustituidas, por lo que pueden ser eliminadas.

```

In [11]: asteroid_df = data1[data1['pha'].notna()]
         asteroid_df = asteroid_df.drop(['diameter', 'albedo', 'diameter_sigma'], axis= 'colu

```

```

In [12]: asteroid_df = asteroid_df[asteroid_df['H'].notna()]
         asteroid_df = asteroid_df[asteroid_df['sigma_ad'].notna()]
         asteroid_df = asteroid_df[asteroid_df['ma'].notna()]

```

```

In [13]: asteroid_df.isnull().sum()

```

```

Out[13]: spkid      0
         neo        0
         pha        0
         H          0
         orbit_id   0
         epoch      0
         epoch_mjd   0
         epoch_cal   0
         e          0
         a          0
         q          0
         i          0
         om         0
         w          0
         ma         0
         ad         0
         n          0
         tp         0
         tp_cal     0
         per        0
         per_y      0
         moid       0
         moid_ld    0
         sigma_e    0
         sigma_a    0
         sigma_q    0
         sigma_i    0
         sigma_om   0
         sigma_w    0
         sigma_ma   0
         sigma_ad   0
         sigma_n    0
         sigma_tp   0
         sigma_per   0
         class      0
         rms        0
         dtype: int64

```

Como podemos ver, ya no tenemos valores nulos en nuestro dataframe.

Algunas columnas presentan valores que no pueden ser procesados dentro de un modelo de machine learning. Estas variables tienen que ser convertidas a variables categóricas.

```

In [14]: asteroid_df['neo'] = asteroid_df['neo'].astype('category')
         asteroid_df['pha'] = asteroid_df['pha'].astype('category')

```

```
asteroid_df['class'] = asteroid_df['class'].astype('category')
```

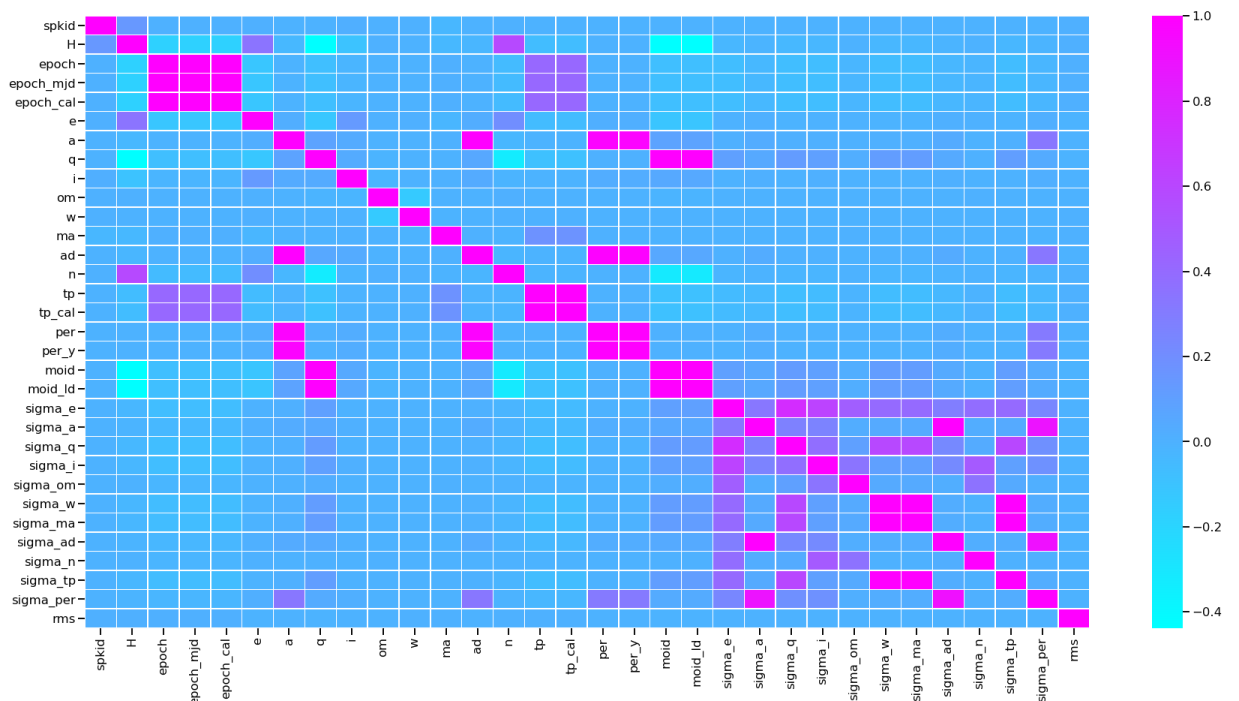
```
In [15]: orbits = asteroid_df['orbit_id'].value_counts().loc[lambdax: x<10].index.to_list()
```

```
In [16]: asteroid_df.loc[asteroid_df['orbit_id'].isin(orbits), 'orbit_id'] = 'other'
```

Observemos la correlación entre las distintas variables.

```
In [61]: plt.figure(figsize = (30, 15))
sns.heatmap(asteroid_df.corr(), annot = False, linewidths=.5, cmap = plt.cm.cool)
```

Out[61]: <AxesSubplot:>



Antes de hacer pasar nuestro dataframe por el modelo, es necesario escalar las variables numéricas. Para ello utilizaremos la función MinMaxScaler.

```
In [17]: asteroid_df = asteroid_df.reset_index(drop=True)#Resetear el índice
```

```
In [18]: #creamos un subset con las variables numericas
subset_df = asteroid_df[asteroid_df.columns[~asteroid_df.columns.isin(['spkid', 'ful
```

```
In [19]: scaler = MinMaxScaler()
scaled_df = scaler.fit_transform(subset_df)
scaled_df = pd.DataFrame(scaled_df, columns = subset_df.columns)
asteroid_df = pd.concat([asteroid_df[['spkid', 'neo', 'pha', 'orbit_id', 'class']],
scaled_df.head())
```

```
Out[19]:
```

	H	epoch	epoch_mjd	epoch_cal	e	a	q	i	om
0	0.131195	0.988218	0.988218	0.989134	0.076017	0.000066	0.030975	0.060467	0.223071
1	0.154519	1.000000	1.000000	1.000000	0.230004	0.000066	0.025712	0.198916	0.480625
2	0.187464	1.000000	1.000000	1.000000	0.256972	0.000063	0.023805	0.074158	0.471810

	H	epoch	epoch_mjd	epoch_cal	e	a	q	i	om	
3	0.119534	0.988218	0.988218	0.989134	0.088732	0.000054	0.025911	0.040748	0.288363	0.416
4	0.233236	1.000000	1.000000	1.000000	0.190939	0.000060	0.025049	0.030614	0.393253	0.996

5 rows × 31 columns




Por último, necesitamos transformar las variables categóricas en variables codificadas. Para ello usaremos la función `get_dummies`

```
In [20]: asteroid_df1 = pd.get_dummies(asteroid_df, columns = ['neo', 'class', 'orbit_id'])
         asteroid_df1.head()
```

```
Out[20]:
```

	spkid	pha	H	epoch	epoch_mjd	epoch_cal	e	a	q	i
0	2000001	N	0.131195	0.988218	0.988218	0.989134	0.076017	0.000066	0.030975	0.060467
1	2000002	N	0.154519	1.000000	1.000000	1.000000	0.230004	0.000066	0.025712	0.198916
2	2000003	N	0.187464	1.000000	1.000000	1.000000	0.256972	0.000063	0.023805	0.074158
3	2000004	N	0.119534	0.988218	0.988218	0.989134	0.088732	0.000054	0.025911	0.040748
4	2000005	N	0.233236	1.000000	1.000000	1.000000	0.190939	0.000060	0.025049	0.030614

5 rows × 242 columns




```
In [21]: lb = LabelBinarizer()
         asteroid_df1['pha'] = lb.fit_transform(asteroid_df1['pha'])
```

```
In [22]: asteroid_df1.head()
```

```
Out[22]:
```

	spkid	pha	H	epoch	epoch_mjd	epoch_cal	e	a	q	i
0	2000001	0	0.131195	0.988218	0.988218	0.989134	0.076017	0.000066	0.030975	0.060467
1	2000002	0	0.154519	1.000000	1.000000	1.000000	0.230004	0.000066	0.025712	0.198916
2	2000003	0	0.187464	1.000000	1.000000	1.000000	0.256972	0.000063	0.023805	0.074158
3	2000004	0	0.119534	0.988218	0.988218	0.989134	0.088732	0.000054	0.025911	0.040748
4	2000005	0	0.233236	1.000000	1.000000	1.000000	0.190939	0.000060	0.025049	0.030614

5 rows × 242 columns



```
In [23]: outputfile = 'asteroid_processed.csv'
         asteroid_df1.to_csv(outputfile, index=False)
```

3. DIVIDIR LOS DATOS

Antes de preparar nuestros modelos de clasificación necesitamos dividir nuestros datos en los sets de entrenamiento y prueba. Como nuestro dataframe presenta datos muy sesgados para nuestra variable objetivo, vamos a utilizar la función StratifiedShuffleSplit para mantener la misma proporción de clases.

```
In [24]: path = 'C:/Users/NITROPC/Desktop/DATA SCIENCE/CERTIFICACION MACHINE LEARNING/06 - PR
          asteroid_processed = pd.read_csv(path, sep = ',')
```

```
In [25]: feature_cols = list(asteroid_processed.columns)
          feature_cols.remove('pha')
```

```
In [26]: X_data = asteroid_processed.drop(['spkid', 'pha'], axis = 1)
          y_data = asteroid_processed['pha']
```

```
In [27]: X_train, X_test, y_train, y_test = train_test_split(X_data, y_data, test_size = 0.3,
```

```
In [28]: X_test.shape
```

```
Out[28]: (279701, 240)
```

```
In [29]: y_test.shape
```

```
Out[29]: (279701,)
```

```
In [30]: X_train.shape
```

```
Out[30]: (652634, 240)
```

```
In [31]: y_train.shape
```

```
Out[31]: (652634,)
```

Dado que nuestros datos presentan un sesgo muy fuerte hacia el No en la variable 'pha' necesitamos que nuestro modelo no sobremuestre la clase positiva. Utilizaremos para ello la librería SMOTE.

```
In [32]: from imblearn.over_sampling import SMOTE
          smtn = SMOTE(random_state = 12)
          X_train_res, y_train_res = smtn.fit_resample(X_train, y_train)
```

```
In [33]: print("Before OverSampling, counts of label 'N': {}".format(sum(y_train == 0)))
          print("Before OverSampling, counts of label 'Y': {} \n".format(sum(y_train == 1)))

          print("After OverSampling, counts of label 'N': {}".format(sum(y_train_res == 0)))
          print("After OverSampling, counts of label 'Y': {}".format(sum(y_train_res == 1)))
```

Before OverSampling, counts of label 'N': 651221
 Before OverSampling, counts of label 'Y': 1413

After OverSampling, counts of label 'N': 651221
 After OverSampling, counts of label 'Y': 651221

4. REGRESIÓN LOGÍSTICA

Vamos a comenzar nuestros modelos con el modelado por regresión logística.

```
In [34]: #creamos el dataframe para nuestras métricas
metrics = pd.DataFrame()
```

```
In [35]: #REgresión Logística estandar
lr = LogisticRegression().fit(X_train_res, y_train_res)
y_pred_lr = lr.predict(X_test)
```

Para el modelo vamos a mostrar las métricas asociadas y la matriz de confusión.

```
In [36]: precision_lr, recall_lr = (round(float(x),2) for x in list(score(y_test,
                                                                    y_pred_lr,
                                                                    average='weighte

# adding lr stats to metrics DataFrame
lr_stats = pd.Series({'precision':precision_lr,
                    'recall':recall_lr,
                    'accuracy':round(accuracy_score(y_test, y_pred_lr), 2),
                    'f1score':round(f1_score(y_test, y_pred_lr), 2),
                    'auc': round(roc_auc_score(y_test, y_pred_lr),2)},
                    name='Logistic Regression')

# Report outcomes
pd.DataFrame(classification_report(y_test, y_pred_lr, output_dict=True)).iloc[:3,:2]
```

```
Out[36]:
```

	0	1
precision	0.999931	0.208279
recall	0.991363	0.970904
f1-score	0.995629	0.342981

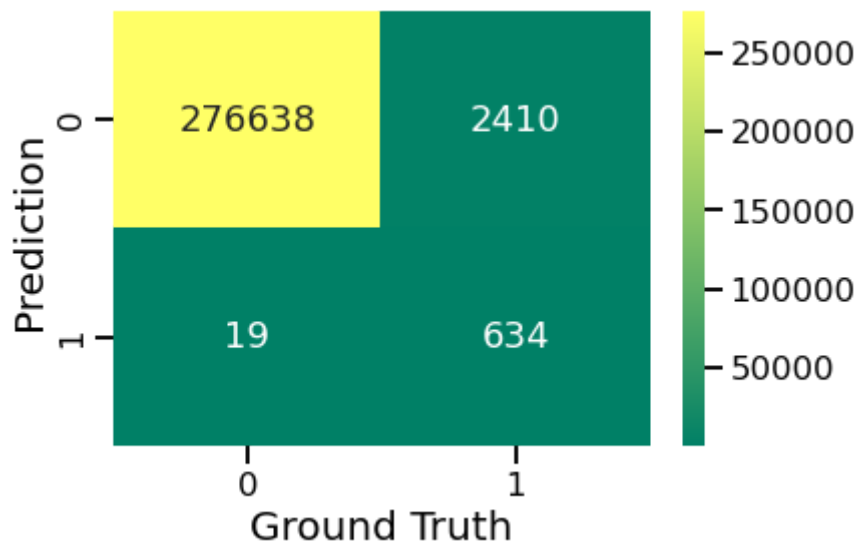
```
In [37]: lr_stats
```

```
Out[37]: precision    1.00
recall      0.99
accuracy    0.99
f1score     0.34
auc         0.98
Name: Logistic Regression, dtype: float64
```

```
In [38]: sns.set_context('talk')
cm = confusion_matrix(y_test, y_pred_lr)
ax = sns.heatmap(cm, annot=True, fmt='d', cmap='summer')

ax.set_ylabel('Prediction', fontsize=20)
ax.set_xlabel('Ground Truth', fontsize=20)
```

```
Out[38]: Text(0.5, 4.5, 'Ground Truth')
```



5. RANDOM FOREST

Nuestro siguiente modelo será Random Forest. Este modelo elimina parte de la posibilidad de sobreajuste.

Podríamos analizar el número de árboles iterando entre varios valores y posteriormente graficar el error para conocer con qué valor se estabiliza el error:

```
rf = RandomForestClassifier(oob_score = True, warm_start = True,
n_jobs = -1, random_state = 1551)

oob_list = list()

for n_trees in [15, 20, 30, 40, 50, 100, 150, 200, 300, 400]:

    rf.set_params(n_estimators = n_trees)
    rf.fit(X_train_res, y_train_res)

    oob_error = 1 - rf.oob_score_
    oob_list.append(pd.Series({'n_trees' : n_trees, 'oob' :
oob_error}))

rf_oob_df = pd.concat(oob_list, axis = 1).T.set_index('n_trees')

sns.set_context('talk')
sns.set_style('white')

ax = rf_oob_df.plot(legend=False, marker='o', color="green", figsize=
(14, 7), linewidth=4)
ax.set(ylabel='out-of-bag error');
```

Sin embargo, para ahorrar tiempo de procesamiento, vamos a aplicar un valor de 150 a `n_estimators`.

```
In [39]: rf = RandomForestClassifier(n_estimators = 150, oob_score = True, warm_start = True,
```

```
In [40]: rf.fit(X_train_res, y_train_res)
```

```
RandomForestClassifier(n_estimators=150, n_jobs=-1, oob_score=True,
```

Out[40]: random_state=1551, warm_start=True)

In [41]: y_pred_rf = rf.predict(X_test)

In [42]:

```
precision_rf, recall_rf = (round(float(x),2) for x in list(score(y_test,
                                                                y_pred_rf,
                                                                average='weighte

rf_stats = pd.Series({'precision':precision_rf,
                    'recall':recall_rf,
                    'accuracy':round(accuracy_score(y_test, y_pred_rf), 2),
                    'f1score':round(f1_score(y_test, y_pred_rf), 2),
                    'auc': round(roc_auc_score(y_test, y_pred_rf),2)}), name='Rando

pd.DataFrame(classification_report(y_test, y_pred_rf, output_dict=True)).iloc[:3,:2]
```

Out[42]:

	0	1
precision	0.999968	0.947059
recall	0.999871	0.986217
f1-score	0.999919	0.966242

In [43]: rf_stats

Out[43]:


```
precision    1.00
recall       1.00
accuracy     1.00
f1score      0.97
auc          0.99
Name: Random Forest, dtype: float64
```

In [44]:

```
sns.set_context('talk')
cm_rf = confusion_matrix(y_test, y_pred_rf)
ax = sns.heatmap(cm_rf, annot=True, fmt='d', cmap='summer')

ax.set_ylabel('Prediction', fontsize=20)
ax.set_xlabel('Ground Truth', fontsize=20)
```

Out[44]: Text(0.5, 4.5, 'Ground Truth')



6.POTENCIADOR DE GRADIENTE

Por último vamos a aplicar un modelo de potenciador de gradiente. Esta técnica combina los principios de la potenciación de gradiente con la aleatoriedad de los árboles de decisión.

```
In [48]: from sklearn.ensemble import GradientBoostingClassifier
```

Al igual que en el modelo Random Forest podríamos analizar el número de árboles iterando entre varios valores y posteriormente graficar el error para conocer con qué valor se estabiliza:

```
error_list = list()

tree_list = [15, 25, 50, 100, 200, 400]
for n_trees in tree_list:

    GBC = GradientBoostingClassifier(n_estimators=n_trees,
    random_state=42)

    print(f'Fitting model with {n_trees} trees')
    GBC.fit(X_train_res.values, y_train_res.values)
    y_pred = GBC.predict(X_test)

    error = 1.0 - accuracy_score(y_test, y_pred)
    Store it
    error_list.append(pd.Series({'n_trees': n_trees, 'error': error}))

error_df = pd.concat(error_list, axis=1).T.set_index('n_trees')

error_df
```

Para ahorrar tiempo de procesamiento, vamos a aplicar un valor de 150 a n_estimators.

```
In [49]: gbc = GradientBoostingClassifier(n_estimators = 400, learning_rate = 0.1, subsample
```

```
In [50]: gbc.fit(X_train_res, y_train_res)
```

```
Out[50]: GradientBoostingClassifier(max_features=4, n_estimators=400, subsample=0.5)
```

```
In [51]: y_pred_gbc = gbc.predict(X_test)
```

```
In [57]: precision_gbc, recall_gbc = (round(float(x),2) for x in list(score(y_test,
                                                                    y_pred_rf,
                                                                    average='weighte

gbc_stats = pd.Series({'precision':precision_gbc,
                        'recall':recall_rf,
                        'accuracy':round(accuracy_score(y_test, y_pred_gbc), 2),
                        'f1score':round(f1_score(y_test, y_pred_gbc), 2),
                        'auc': round(roc_auc_score(y_test, y_pred_gbc),2)}), name='Rand

pd.DataFrame(classification_report(y_test, y_pred_gbc, output_dict=True)).iloc[:3,:2
```

```
Out[57]:
```

	0	1
precision	0.999989	0.705755
recall	0.999029	0.995406
f1-score	0.999509	0.825921


```
In [58]: gbc_stats
```

```
Out[58]: precision    1.00
recall      1.00
accuracy    1.00
f1score     0.83
auc         1.00
Name: Random Forest, dtype: float64
```

```
In [55]: sns.set_context('talk')
cm_rf = confusion_matrix(y_test, y_pred_gbc)
ax = sns.heatmap(cm_rf, annot=True, fmt='d', cmap='summer')

ax.set_ylabel('Prediction', fontsize=20)
ax.set_xlabel('Ground Truth', fontsize=20)
```

Out[55]: Text(0.5, 4.5, 'Ground Truth')



7. CONCLUSIONES

A continuación se resumen las métricas para cada uno de los modelos aplicados:

```
In [59]: metrics.append([lr_stats, rf_stats, gbc_stats])
```

```
Out[59]:
```

	precision	recall	accuracy	f1score	auc
Logistic Regression	1.0	0.99	0.99	0.34	0.98
Random Forest	1.0	1.00	1.00	0.83	1.00
Random Forest	1.0	1.00	1.00	0.83	1.00

Todos los modelos presentan unos valores similares. Sin embargo, dada la naturaleza sesgada

de los datos, se hace necesario el uso de algoritmos de potenciamiento y validación cruzada, que suponen un mayor tiempo de procesado.